

# ARTIFICIAL INTELLIGENCE - UE15CS325

## END SEMESTER ASSESSMENT

---

**VARUN KUMAR S - 01FB15ECS 338**

**VARUN V - 01FB15ECS 341**

**VARUN Y VORA - 01FB15ECS 342**

**VISHWAS S - 01FB15ECS 355**

---

## IMAGE CLASSIFICATION AND LOCALISATION FOR PASCAL VOC 2010

1) DATASET Preparation The size of the dataset is 4474. It was randomly shuffled and was divided as follows:  
Training Size : 2500 Validation Size : 1000 Evaluation Size : 1244

Each image was preprocessed using 'pillow' to reduce its size and padded with 0's This made the dimension of each image (128, 128, 3) The bounding boxes were also modified accordingly

The 15 classes were converted to one-hot vectors. All listed metrics are for top prediction only.

### 2) Building CLASSIFIER

2 convolution layers along with pooling was used. 2 Dense layers were stacked on top of this. A dropout was added for all dense layers. This was followed by a softmax layer. We directly take the highest value as our prediction.

Accuracy : 35.93%

### 3) Object LOCALISATION

2 convolution layers along with pooling was used. 3 Dense Layers with relu activation are stacked on top. The final dense layer gives 4 predictions- xmin, ymin, xmax, ymax.

Accuracy : 63.86%

### 4) Classifier using Transfer Learning

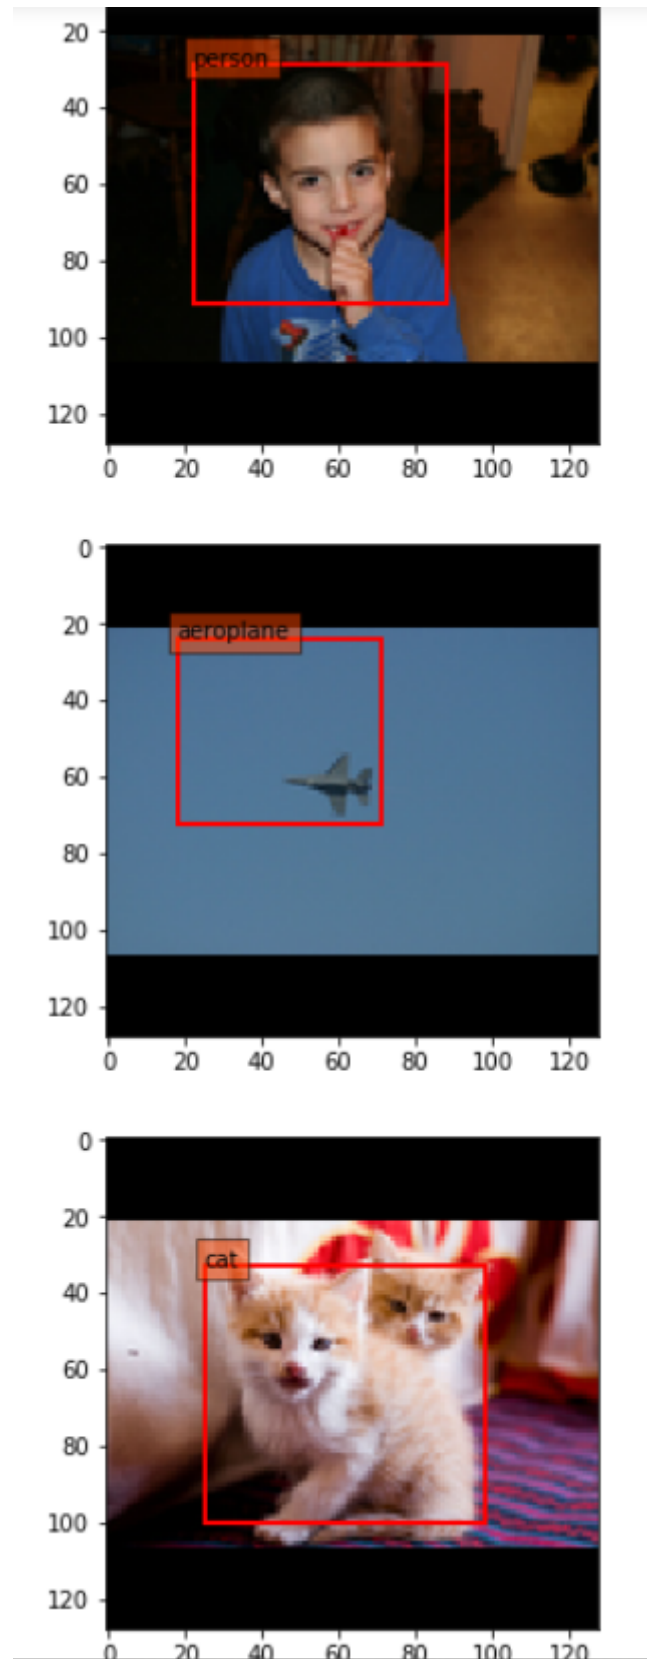
VGG16 model with weights trained in Imagenet was used. We removed the top 2 layers and added 2 Dense Layers with Relu and softmax activations.

Accuracy : 58.5 %

### 5) Localisation using Transfer Learning

VGG16 model was used without its top 2 layers. 3 Dense Layers with 0.5 Dropout. However, a large gain in performance was not observed. Accuracy : 65.54%

## Object Detection



In [18]:

```
from keras.models import load_model

classifier = load_model('classifier_model.h5')
print('\nCLASSIFICATION\n')
classifier.summary()

localisation = load_model('box_model.h5')
print('\nLOCALISAT
```

In [18]:

```
ION\n')
localisation.summary()

classifier_transfer = load_model('transfer_learning_model.h5')
print('\nCLASSIFICATION USING TRANSFER LEARNING\n')
classifier_transfer.summary()

localisation_transfer = load_model('')
print('\nLOCALISATION USING TRANSFER LEARNING\n')
localisation_transfer.summary()
```

CLASSIFICATION

Layer (type)	Output Shape	Param #
=====		
conv2d_9 (Conv2D)	(None, 126, 126, 64)	1792
max_pooling2d_9 (MaxPooling2	(None, 31, 31, 64)	0
conv2d_10 (Conv2D)	(None, 29, 29, 128)	73856
max_pooling2d_10 (MaxPooling	(None, 7, 7, 128)	0
flatten_5 (Flatten)	(None, 6272)	0
dropout_12 (Dropout)	(None, 6272)	0
dense_13 (Dense)	(None, 128)	802944
dropout_13 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 64)	8256
dense_15 (Dense)	(None, 15)	975
=====		
Total params: 887,823		
Trainable params: 887,823		
Non-trainable params: 0		

LOCALISATION

Layer (type)	Output Shape	Param #
=====		
conv2d_5 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_5 (MaxPooling2	(None, 31, 31, 32)	0
conv2d_6 (Conv2D)	(None, 29, 29, 64)	18496
max_pooling2d_6 (MaxPooling2	(None, 7, 7, 64)	0
flatten_3 (Flatten)	(None, 3136)	0
dropout_7 (Dropout)	(None, 3136)	0
=====		

dense_7 (Dense)	(None, 64)	20064
dropout_8 (Dropout)	(None, 64)	0
dense_8 (Dense)	(None, 32)	2080
dropout_9 (Dropout)	(None, 32)	0
dense_9 (Dense)	(None, 4)	132
=====		
Total params: 222,372		
Trainable params: 222,372		
Non-trainable params: 0		

## CLASSIFICATION USING TRANSFER LEARNING

Layer (type)	Output Shape	Param #
=====		
input_4 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
global_average_pooling2d_4 (	(None, 512)	0
dense_7 (Dense)	(None, 128)	65664

---

dense_8 (Dense)	(None, 15)	1935
-----------------	------------	------

---

=====  
Total params: 14,782,287  
Trainable params: 67,599  
Non-trainable params: 14,714,688  
=====

---

In [6]:

```
from NNfunctions import get_dataset, make_confusion_matrix  
x, y, z = get_dataset('C:/Users/Varun/Desktop/AI-ESA/preprocessing')
```

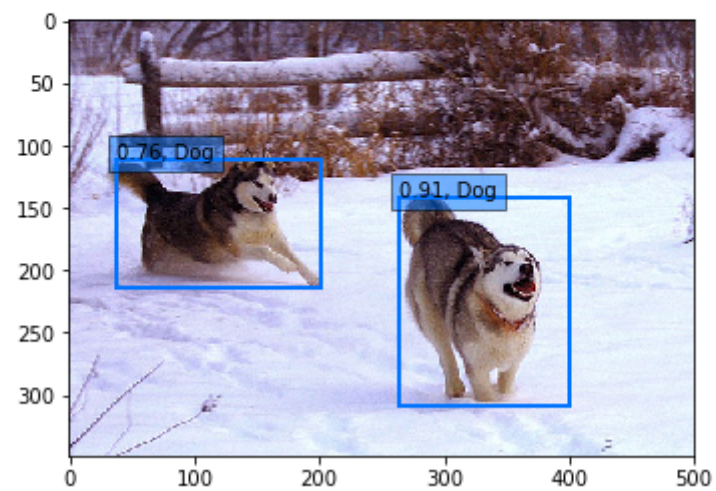

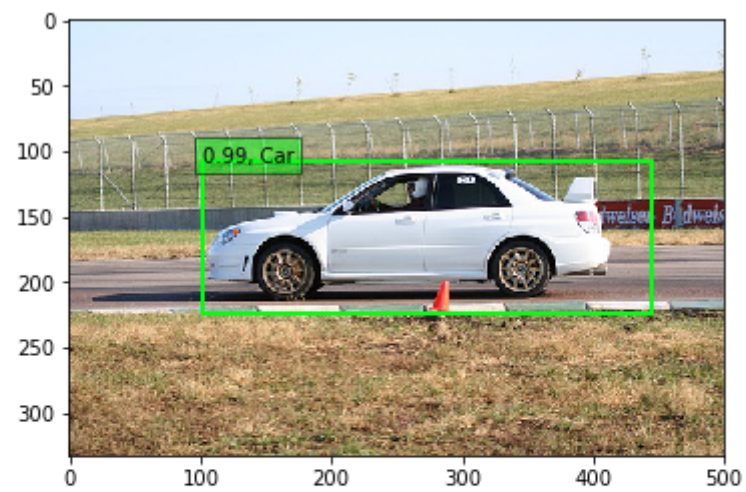
[4744, 4744, 4744]

## Object Detection and classification using TRANSFER LEARNING

Layer (type)	Output Shape	Param #	Connected to
=====			
input_2 (InputLayer)	(None, 300, 300, 3)	0	
conv1_1 (Conv2D)	(None, 300, 300, 64)	1792	input_2[0][0]
conv1_2 (Conv2D)	(None, 300, 300, 64)	36928	conv1_1[0][0]
pool1 (MaxPooling2D)	(None, 150, 150, 64)	0	conv1_2[0][0]
conv2_1 (Conv2D)	(None, 150, 150, 128)	73856	pool1[0][0]
conv2_2 (Conv2D)	(None, 150, 150, 128)	147584	conv2_1[0][0]
pool2 (MaxPooling2D)	(None, 75, 75, 128)	0	conv2_2[0][0]
conv3_1 (Conv2D)	(None, 75, 75, 256)	295168	pool2[0][0]
conv3_2 (Conv2D)	(None, 75, 75, 256)	590080	conv3_1[0][0]
conv3_3 (Conv2D)	(None, 75, 75, 256)	590080	conv3_2[0][0]
pool3 (MaxPooling2D)	(None, 38, 38, 256)	0	conv3_3[0][0]
conv4_1 (Conv2D)	(None, 38, 38, 512)	1180160	pool3[0][0]
conv4_2 (Conv2D)	(None, 38, 38, 512)	2359808	conv4_1[0][0]
conv4_3 (Conv2D)	(None, 38, 38, 512)	2359808	conv4_2[0][0]
pool4 (MaxPooling2D)	(None, 19, 19, 512)	0	conv4_3[0][0]
conv5_1 (Conv2D)	(None, 19, 19, 512)	2359808	pool4[0][0]
conv5_2 (Conv2D)	(None, 19, 19, 512)	2359808	conv5_1[0][0]
conv5_3 (Conv2D)	(None, 19, 19, 512)	2359808	conv5_2[0][0]
pool5 (MaxPooling2D)	(None, 19, 19, 512)	0	conv5_3[0][0]
fc6 (Conv2D)	(None, 19, 19, 1024)	4719616	pool5[0][0]
fc7 (Conv2D)	(None, 19, 19, 1024)	1049600	fc6[0][0]
conv6_1 (Conv2D)	(None, 19, 19, 256)	262400	fc7[0][0]
conv6_2 (Conv2D)	(None, 10, 10, 512)	1180160	conv6_1[0][0]
conv7_1 (Conv2D)	(None, 10, 10, 128)	65664	conv6_2[0][0]
conv7_1z (ZeroPadding2D)	(None, 12, 12, 128)	0	conv7_1[0][0]
conv7_2 (Conv2D)	(None, 5, 5, 256)	295168	conv7_1z[0][0]
conv8_1 (Conv2D)	(None, 5, 5, 128)	32896	conv7_2[0][0]

---

conv4_3_norm_mbox_loc_flat (Flatten)	(None, 17328)	0	conv4_3_norm_mbox_loc[0][0]
fc7_mbox_loc_flat (Flatten)	(None, 8664)	0	fc7_mbox_loc[0][0]
conv6_2_mbox_loc_flat (Flatten)	(None, 2400)	0	conv6_2_mbox_loc[0][0]
conv7_2_mbox_loc_flat (Flatten)	(None, 600)	0	conv7_2_mbox_loc[0][0]
conv8_2_mbox_loc_flat (Flatten)	(None, 216)	0	conv8_2_mbox_loc[0][0]
pool6_mbox_loc_flat (Dense)	(None, 24)	6168	pool6[0][0]
mbox_conf (Concatenate)	(None, 153468)	0	conv4_3_norm_mbox_conf_flat[0][0] fc7_mbox_conf_flat[0][0] conv6_2_mbox_conf_flat[0][0] conv7_2_mbox_conf_flat[0][0] conv8_2_mbox_conf_flat[0][0] pool6_mbox_conf_flat[0][0]
pool6_resaped (Reshape)	(None, 1, 1, 256)	0	pool6[0][0]
mbox_loc (Concatenate)	(None, 29232)	0	conv4_3_norm_mbox_loc_flat[0][0] fc7_mbox_loc_flat[0][0] conv6_2_mbox_loc_flat[0][0] conv7_2_mbox_loc_flat[0][0] conv8_2_mbox_loc_flat[0][0] pool6_mbox_loc_flat[0][0]
mbox_conf_logits (Reshape)	(None, 7308, 21)	0	mbox_conf[0][0]
conv4_3_norm_mbox_priorbox (PriorBox)	(None, 4332, 8)	0	conv4_3_norm[0][0]
fc7_mbox_priorbox (PriorBox)	(None, 2166, 8)	0	fc7[0][0]
conv6_2_mbox_priorbox (PriorBox)	(None, 600, 8)	0	conv6_2[0][0]
conv7_2_mbox_priorbox (PriorBox)	(None, 150, 8)	0	conv7_2[0][0]
conv8_2_mbox_priorbox (PriorBox)	(None, 54, 8)	0	conv8_2[0][0]
pool6_mbox_priorbox (PriorBox)	(None, 6, 8)	0	pool6_resaped[0][0]
mbox_loc_final (Reshape)	(None, 7308, 4)	0	mbox_loc[0][0]
mbox_conf_final (Activation)	(None, 7308, 21)	0	mbox_conf_logits[0][0]
mbox_priorbox (Concatenate)	(None, 7308, 8)	0	conv4_3_norm_mbox_priorbox[0][0] fc7_mbox_priorbox[0][0] conv6_2_mbox_priorbox[0][0] conv7_2_mbox_priorbox[0][0] conv8_2_mbox_priorbox[0][0] pool6_mbox_priorbox[0][0]
predictions (Concatenate)	(None, 7308, 33)	0	mbox_loc_final[0][0] mbox_conf_final[0][0] mbox_priorbox[0][0]





In [17]:

```
from random import randint
k = randint(0,3744)
print('Confusion Matrix for Classification ')
make_confusion_matrix(classifier, x[k:k+500], z[k:k+500])
```

Confusion Matrix for Classification

CONFUSION MATRIX

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14]
[[59  0  2  1  1  0  0 10 12  0  0  1  0  0  0]
 [ 4  3  0  1  0  0  0  1  0  0  7  0  0  0  0]
 [19  0 16  0  7  0  0  2  8  0  1  0  1  0  0]
 [10  0  0 11  6  0  0  1  0  1  2  0  0  0  0]
 [ 0  0  2  5 22  0  0  0  1  1  0  0  0  0  0]
 [ 4  0  0  3  0  0  0  1  2  0  0  0  0  0  0]
 [ 5  0  1  0  0  0  0  2  3  0  0  0  0  0  0]
 [22  0  0  0  1  0  0 43 16  1  1  0  0  0  0]
 [22  0  2  1  2  0  0 14 32  0  1  1  0  0  0]
 [ 1  0  0  3 13  0  0  0  0  1  2  0  0  0  0]
 [11  0  2  7  4  0  0  2  0  0  9  0  0  0  0]
 [ 5  0  1  3  0  0  0  0  0  0  0  1  0  0  0]
 [ 8  0  0  2  0  0  0  0  0  0  0  1  3  0  0]
 [ 4  0  3  1  0  0  0  2  2  0  0  0  0  2  0]
 [ 8  0  0  0  0  0  0  1  0  0  0  0  0  0  0]]
```

{0: 'person', 1: 'tvmonitor', 2: 'bird', 3: 'train', 4: 'aeroplane', 5: 'horse', 6: 'cow', 7: 'cat', 8: 'dog', 9: 'boat', 10: 'car', 11: 'bicycle', 12: 'motorbike', 13: 'bottle', 14: 'chair'}

```
[['ACCURACY' 'PRECISION' 'RECALL']
```

```
['0.3242' '0.6860' '0.2921']
['1.0000' '0.1875' '0.0149']
['0.5517' '0.2963' '0.0792']
['0.2895' '0.3548' '0.0545']
['0.3929' '0.7097' '0.1089']
['0.0000' '0.0000' '0.0000']
['0.0000' '0.0000' '0.0000']
['0.5443' '0.5119' '0.2129']
['0.4211' '0.4267' '0.1584']
['0.2500' '0.0500' '0.0050']
['0.3913' '0.2571' '0.0446']
['0.2500' '0.1000' '0.0050']
['0.7500' '0.2143' '0.0149']
['1.0000' '0.1429' '0.0099']
['0.0000' '0.0000' '0.0000']]
```

In [14]:

```
from random import randint
k = randint(0,3744)
make_confusion_matrix(classifer_transfer, x[k:k+500], z[k:k+500])
#Confusion matrix for a subset of the dataset
```

CONFUSION MATRIX

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14]
[[74  1  1  4  3  2  0  0  1  0  0  0  0  0  1]
 [ 6 14  0  1  0  0  0  0  0  0  0  0  0  0  0]
 [ 7  0 43  1  3  0  0  9  8  0  0  0  0  0  1]
 [ 0  1  0 22  0  0  0  0  0  0  0  0  1  0  0]
 [ 1  0  1  6 30  0  0  1  2  0  1  0  0  0  0]
 [ 4  0  0  2  1  6  0  5  1  0  0  0  0  0  0]
 [ 2  0  0  0  0  0  0  1  7  0  0  0  0  0  0]
 [ 7  2  3  0  0  0  0 47 17  0  0  0  0  0  1]
 [ 9  0  5  1  1  0  0  7 44  0  0  0  0  0  0]
 [ 0  0  0  9  4  0  0  0  1  0  0  0  0  0  0]
 [ 1  0  1  4  4  0  0  0  2  0  9  0  0  0  0]
 [ 3  0  0  3  1  0  0  1  0  0  0  5  1  0  0]
 [ 4  0  0  0  0  0  0  0  1  0  1  2  4  0  0]
 [ 4  0  0  1  0  0  0  2  0  0  0  0  0  4  0]
 [ 3  0  1  2  0  0  0  0  0  0  0  0  0  1  2]]
{0: 'person', 1: 'tvmonitor', 2: 'bird', 3: 'train', 4: 'aeroplane', 5: 'horse', 6: 'cow', 7: 'cat', 8: 'dog', 9: 'boat', 10: 'car', 11: 'bicycle', 12: 'motorbike', 13: 'bottle', 14: 'chair'}
[['ACCURACY' 'PRECISION' 'RECALL']
 ['0.5920' '0.8506' '0.2434']
 ['0.7778' '0.6667' '0.0461']
 ['0.7818' '0.5972' '0.1414']
 ['0.3929' '0.9167' '0.0724']
 ['0.6383' '0.7143' '0.0987']
 ['0.7500' '0.3158' '0.0197']
 ['0.0000' '0.0000' '0.0000']
 ['0.6438' '0.6104' '0.1546']
 ['0.5238' '0.6567' '0.1447']
 ['0.0000' '0.0000' '0.0000']
 ['0.8182' '0.4286' '0.0296']
 ['0.7143' '0.3571' '0.0164']
 ['0.6667' '0.3333' '0.0132']
 ['0.8000' '0.3636' '0.0132']
 ['0.4000' '0.2222' '0.0066']]
```

## Observations and Conclusion:

1. The metrics for car show that the model is able to classify an object as a car relatively well.
2. The most confusion is between cat and dog.
3. The bounded boxes appear to around the same region for most predictions. Probably the model has learnt where the image most likely appears.
4. Shallower models perform better than deeper models when the data is less.
5. Accuracy of models trained using transfer learning outperformed models trained from scratch.

In [ ]:

